

Main Report HW05 - Performance Testing

Student ID: 23127522
SUT: EShop backend, Node/Express + SQLite
Base URL: http://localhost:3000
Tools: Apache JMeter 5.6.3, JMeter Plugins Ultimate Thread Group, Task Manager/Resource Monitor, AI assistant
Main run date: 2026-08-18

1. Summary

This assignment performs performance testing for the EShop backend with JMeter using three main scenarios: Load, Stress, and Spike. All three scenarios use the same end-to-end workflow, called Flow A: the user logs in, reads the profile, browses categories/products, adds an item to the cart, applies a coupon, checks out, and views the order list.

Key results:

Scen ario	Main raw .jtl	HTML report	Raw rows	Max active threads	Total P95	Total P99	Throug hput
Load	reports/results/23127522_Load_20260817.jtl	reports/load	5,000	81	19 ms	30 ms	33.17 req/s
Stress	reports/results/23127522_Stress_20260818.jtl	reports/stress	147,593	500	1,635 ms	5,163.87 ms	290.01 req/s
Spike	reports/results/23127522_Spike_20260818.jtl	reports/spike	32,850	300	448 ms	794 ms	157.21 req/s

The Load test is a clean baseline with 0 errors. The Stress test shows a clear latency increase when the load reaches 500 active threads, but it does not show mass child HTTP failures. The Spike test is valid after reseeding the data; the earlier Spike run is excluded from the final conclusion because it contained many 401/403 errors caused by dirty data/login state.

2. Environment, Data, And Workflow

2.1 Runtime Environment

Component	Value
Hostname	LAPTOP-N9D07817
CPU	Intel(R) Core(TM) Ultra 7 155H
Logical cores	22
RAM	15.71 GB
OS	Microsoft Windows 11 Home Single Language
Backend process	node.exe

Component	Value
JMeter process	java.exe

The backend uses SQLite, and the data can change after each restart/test run. Therefore, before each official run, I execute:

```
node scripts\seed-data.js
node scripts\verify-seed.js
```

Measurement limitation: JMeter and the backend run on the same machine, so the results represent the local hardware threshold of this machine, not production capacity.

2.2 Flow A

Step	Request	Endpoint group	Note
1	POST /api/login	Auth-heavy	Gets the JWT token for later requests.
2	GET /api/users/me	Auth-heavy	Verifies that the token is valid.
3	GET /api/categories	Read-heavy	Reads product categories.
4	GET /api/products/:id	Read-heavy	Reads product detail using data from CSV.
5	POST /api/cart	Transactional	Adds a product to the cart.
6	GET /api/cart	Transactional	Checks cart state.
7	POST /api/apply-coupon	Transactional	Applies a fixed coupon.
8	POST /api/checkout	Transactional	Creates an order.
9	GET /api/orders/my-orders	Transactional	Reads the user's order list.

The requests are placed inside a Transaction Controller to measure the end-to-end time of the entire flow. The login request uses a JSON Extractor to store the token in `${authToken}`. Authenticated requests then use an HTTP Header Manager:

```
Authorization: Bearer ${authToken}
```

2.3 Data-Driven Testing

File	Role
data/users.csv	500 main users for Load/Stress/Spike.
data/products.csv	Product IDs, categories, and prices for data-driven product/cart/checkout requests.
data/coupons.csv	Fixed coupons, avoiding the percent-coupon bug that miscalculates <code>final_amount</code> .
data/users_lockout.csv	Separate users for lockout testing, so the 500 main users are not polluted.

Data-driven testing means the test plan does not hard-code a single user/product. JMeter reads each CSV row, maps values into variables, and uses those variables in requests. This allows each virtual user to follow the same business flow with different data.

3. How To Run And Read JMeter Output

3.1 General Preparation

Start the backend in a separate terminal:

```
cd /d "<project-root>\eshop-sut\backend"
node server.js
```

In the test terminal:

```
cd /d "<project-root>"
node scripts\seed-data.js
node scripts\verify-seed.js
```

When running JMeter in non-GUI mode, remove the old `.jtl` file and the old report directory first. Otherwise, JMeter may report that the file/folder already exists or append to an old result, which would corrupt the numbers.

3.2 Main Output Meaning

Output	Meaning
<code>.jtl</code>	Raw log for each sample. This is the primary source for checking pass/fail, response code, failing label, and active threads.
<code>statistics.json</code>	Aggregated file generated by the HTML dashboard, used for avg, median, p90, p95, p99, and throughput.
<code>index.html</code>	Visual dashboard generated by JMeter from <code>.jtl</code> ; used for presentation and graph review.
Console summary +	Statistics for the latest interval.
Console summary =	Cumulative statistics, but it should not be used alone if the <code>.jtl</code> file contains data.

Common fields in `.jtl`:

Field	Meaning
<code>timeStamp</code>	Time when the sample was recorded, in epoch milliseconds.
<code>elapsed</code>	Total sample processing time in ms.
<code>label</code>	Sampler or transaction name, for example <code>01 POST login</code> or <code>Flow A - Login Browse Cart Checkout</code> .
<code>responseCode</code>	HTTP status code. If it is blank on a transaction parent, it is not a real HTTP response.

Field	Meaning
responseMessage	Message associated with the response code.
success	true / false based on assertions and sampler status.
failureMessage	Failure reason, if any.
bytes	Number of bytes received.
sentBytes	Number of bytes sent.
grpThreads	Active threads in the Thread Group at sample time.
allThreads	Total active threads at sample time.
Latency	Time until the first byte is received, in ms.
Connect	Connection setup time, in ms.

Metrics in `statistics.json` /HTML report:

Metric	Meaning
Samples	Number of samples counted in the dashboard. It may differ from <code>.jt1</code> row count if transaction parents or filters are involved.
Error %	Percentage of failed samples. The raw <code>.jt1</code> must be checked to know whether failures are child HTTP failures or transaction parent artifacts.
Average	Mean response time. It is sensitive to outliers.
Median	P50, meaning 50% of requests are faster than this value.
90% Line / P90	90% of requests are faster than this value.
95% Line / P95	95% of requests are faster than this value; this is the main metric for detecting regression.
99% Line / P99	99% of requests are faster than this value; used for tail latency.
Throughput	Number of requests per second handled during the run.
Received KB/sec	Data receive rate from the server.
Sent KB/sec	Data send rate to the server.

4. Sequential Review Of The Three Test Plans

4.1 Load Test

Goal: measure a stable baseline when 500 virtual users run Flow A once with a moderate ramp-up. This scenario shows how fast/slow the system is under normal expected load before increasing pressure.

Files and report:

Component	Value
Test plan	test-plans/23127522_Load_20260817.jmx
Listener/report view	Aggregate Report
Raw result	reports/results/23127522_Load_20260817.jtl
HTML report	reports/load/index.html
Screenshot	reports/evidence/23127522_Load_20260817.png

How to run:

```
cd /d "<project-root>"

if not exist reports\results mkdir reports\results
rmdir /s /q reports\load
del /q reports\results\23127522_Load_20260817.jtl

jmeter -n ^
-t test-plans\23127522_Load_20260817.jmx ^
-l reports\results\23127522_Load_20260817.jtl ^
-e -o reports\load ^
-Jsummariser.name=summary ^
-Jsummariser.interval=10 ^
-Jsummariser.out=true ^
-Jsummariser.log=true
```

Configuration explanation:

Parameter	Value	Meaning
Number of Threads	500	Total virtual users created during the run.
Ramp-up period	120s	JMeter does not start all 500 users immediately; it spreads them over 120 seconds.
Loop Count	1	Each user runs Flow A once.
Think time	Per plan	Simulates user thinking time between actions.
Aggregate Report	1 listener	Used to review avg, median, percentiles, and throughput per sampler.

Result:

Metric	Value
Raw rows	5,000

Metric	Value
Pass	5,000
Fail	0
Max active threads observed in <code>.jtl</code>	81
Total samples in <code>statistics.json</code>	4,500
Error %	0.0000%
Avg	8.16 ms
Median	6 ms
P90	15 ms
P95	19 ms
P99	30 ms
Throughput	33.17 req/s

Analysis: The Load test is clean because all 5,000 raw rows passed. P95 of 19 ms and P99 of 30 ms show a very low baseline. Max active threads reached only 81 even though the plan configured 500 users because each user completed the flow quickly and ramp-up lasted 120 seconds. Therefore, this Load test measures 500 total users, not 500 concurrent users at every moment.

4.2 Stress Test

Goal: increase the load in steps to observe how latency, throughput, and errors change as active threads approach 500. A Stress test does not only ask whether the system passes; it asks where the system starts to degrade.

Files and report:

Component	Value
Test plan	<code>test-plans/23127522_Stress_20260818.jmx</code>
Listener/report view	Summary Report
Raw result	<code>reports/results/23127522_Stress_20260818.jtl</code>
HTML report	<code>reports/stress/index.html</code>
Screenshot	<code>reports/evidence/23127522_Stress_20260818.png</code>

How to run:

```
cd /d "<project-root>"

if not exist reports\results mkdir reports\results
rmdir /s /q reports\stress
del /q reports\results\23127522_Stress_20260818.jtl

jmeter -n ^
```

```
-t test-plans\23127522_Stress_20260818.jmx ^
-l reports\results\23127522_Stress_20260818.jtl ^
-e -o reports\stress ^
-Jsummariser.name=summary ^
-Jsummariser.interval=10 ^
-Jsummariser.out=true ^
-Jsummariser.log=true
```

Ultimate Thread Group configuration explanation:

Step	Added threads	Initial delay	Startup	Hold	Shutdown	Meaning
1	100	0s	30s	420s	10s	Creates a 100-user baseline.
2	100	90s	30s	330s	10s	Increases to about 200 users.
3	100	180s	30s	240s	10s	Increases to about 300 users.
4	100	270s	30s	150s	10s	Increases to about 400 users.
5	100	360s	30s	60s	10s	Increases to about 500 users.

Parameter meaning:

Parameter	Meaning
threads	Number of virtual users for that step.
initial_delay	How long to wait from test start before this step starts.
startup_time	Ramp-up time for that step.
hold_load_for	How long to hold the load after the step finishes ramping up.
shutdown_time	How long to ramp down users from that step.

Overall result:

Metric	Value
Raw rows	147,593
Pass	147,550
Fail	43
Max active threads	500
Duration	459.36s
Total samples in statistics.json	133,090
Error %	0.0323%
Avg	444.95 ms
Median	907 ms
P90	1,420 ms

Metric	Value
P95	1,635 ms
P99	5,163.87 ms
Throughput	290.01 req/s

Slowest endpoints by P95:

Sampler	Samples	Errors	Avg	P95	P99
Flow A parent	15,003	43	3,868.12 ms	8,370.80 ms	8,763.96 ms
07 POST apply coupon	14,637	0	757.44 ms	1,662 ms	1,857 ms
01 POST login	14,960	0	488.15 ms	1,194 ms	1,375 ms
08 POST checkout	14,578	0	450.00 ms	1,010 ms	1,144.21 ms
09 GET my orders	14,522	0	446.36 ms	1,010 ms	1,146.77 ms

Slice by active threads:

Active threads	Samples	Avg	P95	P99	Approx. throughput
1-100	12,942	39 ms	75 ms	573 ms	28.20 req/s
101-200	25,337	112.16 ms	288 ms	401 ms	69.04 req/s
201-300	30,448	303.21 ms	606 ms	824 ms	110.81 req/s
301-400	32,353	538.04 ms	1,008 ms	1,219 ms	176.16 req/s
401-500	31,510	853.89 ms	1,526 ms	1,763 ms	347.60 req/s

Analysis: The Stress test achieved the goal of reaching 500 active threads. As load increased, latency increased clearly: child HTTP P95 rose from about 606 ms in the 201-300 bucket to 1,526 ms in the 401-500 bucket. The most notable endpoint is 07 POST apply coupon , with P95 of 1,662 ms. The 43 errors are not backend HTTP failures; all of them are on the transaction parent Flow A - Login Browse Cart Checkout , with blank response code and Response was null , while child HTTP requests contain 147,550 200 rows.

4.3 Spike Test

Goal: check how the system reacts when load increases suddenly. Unlike Stress, Spike does not increase load through long gradual steps; it creates a fast jump from a 20 VU baseline to 300 VU.

Files and report:

Component	Value
Test plan	test-plans/23127522_Spike_20260818.jmx
Listener/report view	View Results Tree
Main raw result	reports/results/23127522_Spike_20260818.jtl
Main HTML report	reports/spike/index.html

Component	Value
Screenshot	reports/evidence/23127522_Spike_20260818.png

How to run:

```
cd /d "<project-root>"

if not exist reports\results mkdir reports\results
rmdir /s /q reports\spike
del /q reports\results\23127522_Spike_20260818.jtl

jmeter -n ^
-t test-plans\23127522_Spike_20260818.jmx ^
-l reports\results\23127522_Spike_20260818.jtl ^
-e -o reports\spike ^
-Jsummariser.name=summary ^
-Jsummariser.interval=10 ^
-Jsummariser.out=true ^
-Jsummariser.log=true
```

In this report, I use the valid rerun at reports/results/23127522_Spike_20260818.jtl and reports/spike , because the earlier Spike run had data/login issues.

Ultimate Thread Group configuration explanation:

Phase	Threa ds	Initial delay	Start up	Hol d	Shutdo wn	Meaning
Baseli ne	20	0s	10s	170 s	10s	Keeps low background load to show stable behavior before the spike.
Spike	280	60s	5s	30s	5s	Quickly adds 280 users, reaching a peak of 300 users.

Valid rerun result:

Metric	Value
Raw rows	32,850
Pass	32,809
Fail	41
Max active threads	300
Duration	189.38s
Total samples in statistics.json	29,689
Error %	0.1381%
Avg	105.90 ms
Median	47 ms

Metric	Value
P90	288 ms
P95	448 ms
P99	794 ms
Throughput	157.21 req/s

Excluded invalid Spike run:

File	Rows	200	401	403	Blank	Failed
Earlier invalid Spike run	41,238	12,371	8,444	20,385	38	28,867

Analysis: Spike has P95 of 448 ms and P99 of 794 ms, much better than Stress because the duration is shorter and the peak load is lower than Stress. The earlier failed run is not used for performance conclusions because mass 401/403 errors reflect dirty data/login state. After reseeding and verifying data, the child HTTP requests of the Spike run return 200 ; the remaining 41 failures are blank transaction parent records and must be reviewed in .jtl instead of being treated as backend failures.

5. Monitoring And Evidence

During test execution, I monitor:

Process	Meaning
node.exe	EShop backend, the SUT whose CPU/RAM should be observed.
java.exe	JMeter load generator, which must be separated from backend resource usage.

Evidence:

Scena rio	Screenshot	Raw result	Report
Load	reports/evidence/23127522_Load_20260817.png	reports/results/23127522_Load_20260817.jtl	reports/load/index.html
Stress	reports/evidence/23127522_Stress_20260818.png	reports/results/23127522_Stress_20260818.jtl	reports/stress/index.html
Spike	reports/evidence/23127522_Spike_20260818.png	reports/results/23127522_Spike_20260818.jtl	reports/spike/index.html

Demo video: https://youtu.be/GPZ94lcdWDI?si=GJrB3ISKFIGie_b2

6. AI Preview And Human Review

This section summarizes how I used AI for previews/recommendations, then verified the conclusions against raw .jtl and statistics.json . I do not use AI output as the final conclusion unless it has been checked against real measurements.

6.1 AI Preview

AI helped propose:

AI-supported item	Human review
End-to-end Flow A design.	Kept because it covers auth-heavy, read-heavy, and transactional endpoints.
CSV-based data-driven testing.	Kept, with lockout users separated to avoid locking the main users.
Load/Stress/Spike with different load levels.	Adjusted Stress/Spike to Ultimate Thread Group to model continuous load more correctly.
P95 thresholds and optimization ideas.	Accepted only after checking raw <code>.jtl</code> , <code>statistics.json</code> , and code behavior.

6.2 Misinterpretations Corrected

Risky interpretation	Wrong conclusion	Correct value from raw log	Correction
Console shows <code>summary = 0</code>	The test did not run.	Load <code>.jtl</code> has 5,000 rows, 5,000 pass, 0 fail.	Do not conclude from one console line; check <code>.jtl</code> and report.
Stress has 43 failed samples	Backend returned HTTP errors.	All 43 failures are transaction parent <code>Response was null</code> ; child HTTP has 147,550 <code>200</code> rows.	This is a transaction parent/shutdown artifact, not an HTTP backend failure.
Stress has 147,593 rows	There were 147,593 users.	Max active threads is 500.	<code>.jtl</code> rows are sample records; each user can generate many requests and multiple loops.
Load only shows max active threads 81	Only 81 users ran.	<code>.jmx</code> has 500 threads, ramp-up 120s, loop 1; raw <code>.jtl</code> has 500 Flow A parent samples.	500 is total virtual users, not concurrent users at every moment.
Earlier Spike run has high error rate	Backend fails under spike.	Invalid run has 8,444 <code>401</code> , 20,385 <code>403</code> , failed 28,867/41,238 rows.	This run is invalid due to dirty data/login state; use the valid rerun for analysis.

6.3 Optimization Recommendation Review

Recommendation	Classification	Reason
Add index <code>orders(user_id)</code>	Feasible	<code>GET /api/orders/my-orders</code> filters by user, so an index reduces scans as orders grow.
Add pagination for <code>my-orders</code>	Feasible	Prevents the response from growing indefinitely when a user has many orders.
Clear <code>userCarts</code> after checkout	Feasible	Reduces memory state and prevents old carts from affecting later loops.

Recommendation	Classification	Reason
Enable SQLite WAL	Feasible/needs benchmark	May improve read/write concurrency with SQLite.
Reduce logging in hot paths	Needs consideration	Useful if log I/O is heavy, but should not remove debug capability.
Separate the load generator from the SUT	Needs consideration	Improves measurement cleanliness, but it is not a backend optimization by itself.
Generic connection pool for embedded SQLite	Hallucinated/not suitable	Embedded SQLite is not a client-server DB; a generic pool does not automatically solve write bottlenecks.

7. Task 3 - Continuous Performance Testing Model

7.1 Definition

Continuous Performance Testing is a model that brings performance testing into the continuous development lifecycle. Each commit or pull request is classified by risk, then mapped to the appropriate performance checks. The goal is to detect p95/p99 regression, increased error rate, or reduced throughput before a change is merged.

In this assignment, CPT does not mean running the full Stress/Spike suite on every commit. A tiered strategy is more suitable because it balances reliability and cost.

7.2 Tiered Test Model

Change type	Test to run	Reason
Docs/report/evidence only	Skip performance test	Does not affect SUT runtime.
JMX/data/scripts	Validate JMX + seed + verify + smoke	These changes can easily break data, token handling, or the test plan.
Backend/API/DB	PR Load Gate	Detects early regression in Flow A.
Nightly/release/manual	Full suite: Load + Stress + Spike	Provides trend coverage for baseline, degradation, and spike behavior.
Weekly/manual soak	Endurance 10-15 minutes	Checks memory growth and p95 drift over time.

7.3 Alert Rules


Rule	Result
p95 increases by > 20% compared with baseline and absolute increase is > 100 ms	Flag p95 regression.
Child HTTP error rate > 0.5% in Load/PR gate	Fail build.
Many 5xx, timeout, or connection refused errors	Fail build.

Rule	Result
Mass 401/403 errors	Mark invalid, reseed, and rerun.
Stress/Spike has parent Response was null but child HTTP 200	Warning, requires human .jtl review.
Throughput decreases by > 15% while active threads are comparable	Flag throughput regression.

Proposed hard thresholds:

Scenario	Proposed threshold
Load	Total P95 <= 500 ms, child HTTP error = 0%.
Stress	Total P95 <= 3,000 ms, child HTTP error < 1%, backend must not crash.
Spike	Total P95 <= 1,500 ms, child HTTP error < 1%, throughput should recover after the spike.
Proposed Endurance	250-300 active threads for 10-15 minutes; requires a separate soak run for official confirmation.

7.4 Flowchart



7.5 Trade-Offs

Trade-off	Benefit	Risk/handling
Do not run the full suite on every PR	Saves time	May miss issues that only appear under Stress/Spike; compensate with nightly runs.
Relative p95 threshold of 20%	Detects regression against the local baseline	Can create false alarms when baseline is very small; add absolute increase > 100 ms.
JMeter and backend on the same machine	Easy local setup	Resource usage is mixed; clearly state that this is a local threshold.

Trade-off	Benefit	Risk/handling
Reseed data every run	Avoids fake 401/403 errors	Takes extra time but is necessary.
Manual baseline update	Avoids normalizing regressions	Requires human review before accepting a new baseline.

8. Bugs And Performance Issues

Issue	Type	Impact	GitHub
Lockout actually happens after 2 wrong attempts instead of the expected 3	Functional/security	Medium	#3
Percent coupon calculates <code>final_amount</code> incorrectly	Functional	High if percent coupons are used	#5
<code>apply-coupon</code> does not require authentication	Security/logic	Medium	#2
<code>price</code> sometimes returns as string and sometimes as number	API consistency	Low-Medium	#4
<code>my-orders</code> lacks pagination/index	Performance risk	Medium-High	#6

9. Conclusion

The Load test provides a stable baseline: all 5,000 raw rows passed, with Total P95 of 19 ms and P99 of 30 ms. The Stress test reached 500 active threads and 290.01 req/s throughput; latency increased clearly with Total P95 of 1,635 ms and P99 of 5,163.87 ms, but the 43 recorded failures are transaction parent `Response was null` records, not child HTTP failures. The Spike test reached 300 active threads, with Total P95 of 448 ms and P99 of 794 ms; the earlier Spike run is excluded because it had many 401/403 errors caused by dirty data/login state.

The optimizations to prioritize are adding an index on `orders(user_id)`, adding pagination for `my-orders`, clearing `userCarts` after checkout, and benchmarking SQLite WAL. The proposed CPT model helps detect p95 regression earlier in future commits, but it must always reseed clean data and check raw `.jt1` before drawing conclusions.